

Smart Assistant for Research Summarization

This document is created to test the document upload, extraction, and chunking pipeline.

The goal of this project is to reduce the time required to read large research papers.

Traditional summarization tools fail to provide deep comprehension and reasoning.

This system uses Retrieval Augmented Generation to answer questions.

It also evaluates user answers and provides justification from the source document.

Chunking is a critical step in the RAG pipeline.

Poor chunking leads to poor retrieval and incorrect answers.

FAISS is used as the vector database for semantic search.

FastAPI powers the backend of this application.